

# Interactive Real-Time ML Serving Architecture

From a user-triggered traffic spike to observable autoscaling, inference, and safe rollback.

PORTFOLIO DEMO

